

Homework2DenseNet

February 28, 2022

1 Homework 2 DenseNet121

1.1 Loading Libraries

```
[1]: # Importing Libraries

# Basic Libraries
→#####

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import seaborn as sns
import datetime
import time

# For Feature Engineering
→#####

# For Machine Learning Techniques
→#####

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import cifar10

from tensorflow.keras.applications import DenseNet121
from tensorflow.keras.applications.densenet import preprocess_input

# For Data Analysis
→#####

from sklearn.model_selection import train_test_split
```

```

from sklearn.metrics import classification_report, confusion_matrix

# Personal Preference
→#####
import warnings
warnings.filterwarnings('ignore')

```

```

from tensorflow.keras.preprocessing import image from tensorflow.keras.preprocessing.image import ImageDataGenerator,img_to_array

```

```

from tensorflow.keras.models import Model from tensorflow.keras.optimizers import Adam from tensorflow.keras.callbacks import ModelCheckpoint, ReduceLROnPlateau

```

```

from tensorflow.keras.layers import Dense,GlobalAveragePooling2D,Convolution2D,BatchNormalization
from tensorflow.keras.layers import Flatten,MaxPooling2D,Dropout

```

1.1.1 Setting GPU

[2]: *# Change to markdown if gpu is not set up*

```

import os

os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2"
from tensorflow.python.client import device_lib
print(device_lib.list_local_devices())

physical_devices = tf.config.list_physical_devices("GPU")
tf.config.experimental.set_memory_growth(physical_devices[0], True)

```

```

[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 13338002541111118179
xla_global_id: -1
, name: "/device:GPU:0"
device_type: "GPU"
memory_limit: 6273040384
locality {
  bus_id: 1
  links {
  }
}
incarnation: 5586588163853313758
physical_device_desc: "device: 0, name: NVIDIA GeForce RTX 2070 SUPER, pci bus
id: 0000:02:00.0, compute capability: 7.5"

```

```
xla_global_id: 416903419  
]
```

1.2 Loading Data

```
[3]: # Setting Class Names
```

```
class_names=['airplane','automobile','bird','cat','deer',  
             'dog','frog','horse','ship','truck']
```

```
[4]: # Loading the Dataset
```

```
(x_train,y_train),(x_test,y_test)=cifar10.load_data()
```

```
[5]: # Normalizing the Images
```

```
x_train=x_train/255.0  
print(x_train.shape)
```

```
x_test=x_test/255.0  
print(x_test.shape)
```

```
(50000, 32, 32, 3)
```

```
(10000, 32, 32, 3)
```

1.3 Splitting the Data

```
[6]: # 10% of the Original Dataset
```

```
x_train10, not_needed1, y_train10, not_needed2 = train_test_split(  
    x_train, y_train, test_size=0.90, random_state=42)
```

```
[7]: # 50% of the Original Dataset
```

```
x_train50, not_needed1, y_train50, not_needed2 = train_test_split(  
    x_train, y_train, test_size=0.50, random_state=42)
```

```
[8]: # 80% of the Original Dataset (redundant)
```

```
x_train80, not_needed1, y_train80, not_needed2 = train_test_split(  
    x_train, y_train, test_size=0.20, random_state=42)
```

1.4 Setting up Models

```
[9]: # Creating the actual model
```

```
model_10 = tf.keras.applications.DenseNet121(  
    include_top=True,  
    weights=None,  
    input_tensor=None,  
    input_shape=(32, 32, 3),  
    pooling=None,  
    classes=1000,)
```

```
[10]: # Setting up Mutiple Models
```

```
model_50 = model_10  
  
model_80 = model_10
```

```
[11]: # Model Compiling
```

```
model_10.compile(loss="sparse_categorical_crossentropy",  
                 optimizer="Adam", metrics=["sparse_categorical_accuracy"])  
  
model_50.compile(loss="sparse_categorical_crossentropy",  
                 optimizer="Adam", metrics=["sparse_categorical_accuracy"])  
  
model_80.compile(loss="sparse_categorical_crossentropy",  
                 optimizer="Adam", metrics=["sparse_categorical_accuracy"])
```

1.5 Training the Models

```
[12]: # Fitting the 10% Model
```

```
history10 = model_10.fit(x_train10,y_train10,epochs=10)
```

Epoch 1/10

157/157 [=====] - 25s 71ms/step - loss: 2.1026 -
sparse_categorical_accuracy: 0.2984

Epoch 2/10

157/157 [=====] - 10s 62ms/step - loss: 1.7068 -
sparse_categorical_accuracy: 0.3742

Epoch 3/10

157/157 [=====] - 10s 64ms/step - loss: 1.4992 -
sparse_categorical_accuracy: 0.4578

Epoch 4/10

157/157 [=====] - 10s 62ms/step - loss: 1.4213 -
sparse_categorical_accuracy: 0.4838

```

Epoch 5/10
157/157 [=====] - 10s 62ms/step - loss: 1.2623 -
sparse_categorical_accuracy: 0.5432
Epoch 6/10
157/157 [=====] - 10s 62ms/step - loss: 1.1838 -
sparse_categorical_accuracy: 0.5740
Epoch 7/10
157/157 [=====] - 10s 65ms/step - loss: 1.0485 -
sparse_categorical_accuracy: 0.6296
Epoch 8/10
157/157 [=====] - 10s 63ms/step - loss: 0.9555 -
sparse_categorical_accuracy: 0.6622
Epoch 9/10
157/157 [=====] - 10s 61ms/step - loss: 0.8585 -
sparse_categorical_accuracy: 0.6952
Epoch 10/10
157/157 [=====] - 10s 61ms/step - loss: 0.7603 -
sparse_categorical_accuracy: 0.7364

```

[13]: *# Fitting the 50% Model*

```
history50 = model_50.fit(x_train50,y_train50,epochs=10)
```

```

Epoch 1/10
782/782 [=====] - 49s 63ms/step - loss: 1.1217 -
sparse_categorical_accuracy: 0.6128
Epoch 2/10
782/782 [=====] - 48s 61ms/step - loss: 0.9520 -
sparse_categorical_accuracy: 0.6694
Epoch 3/10
782/782 [=====] - 49s 62ms/step - loss: 0.8130 -
sparse_categorical_accuracy: 0.7200
Epoch 4/10
782/782 [=====] - 49s 62ms/step - loss: 0.7176 -
sparse_categorical_accuracy: 0.7540
Epoch 5/10
782/782 [=====] - 48s 62ms/step - loss: 0.6150 -
sparse_categorical_accuracy: 0.7871
Epoch 6/10
782/782 [=====] - 48s 62ms/step - loss: 0.5415 -
sparse_categorical_accuracy: 0.8129
Epoch 7/10
782/782 [=====] - 48s 61ms/step - loss: 0.4448 -
sparse_categorical_accuracy: 0.8457
Epoch 8/10
782/782 [=====] - 49s 62ms/step - loss: 0.3821 -
sparse_categorical_accuracy: 0.8667
Epoch 9/10

```

```
782/782 [=====] - 49s 63ms/step - loss: 0.3162 -  
sparse_categorical_accuracy: 0.8884  
Epoch 10/10  
782/782 [=====] - 49s 62ms/step - loss: 0.2573 -  
sparse_categorical_accuracy: 0.9107
```

```
[14]: # Fitting the 80% Model
```

```
history80 = model_80.fit(x_train80,y_train80,epochs=10)
```

```
Epoch 1/10  
1250/1250 [=====] - 84s 64ms/step - loss: 0.4941 -  
sparse_categorical_accuracy: 0.8419  
Epoch 2/10  
1250/1250 [=====] - 79s 63ms/step - loss: 0.3693 -  
sparse_categorical_accuracy: 0.8741  
Epoch 3/10  
1250/1250 [=====] - 79s 63ms/step - loss: 0.2946 -  
sparse_categorical_accuracy: 0.8997  
Epoch 4/10  
1250/1250 [=====] - 80s 64ms/step - loss: 0.2465 -  
sparse_categorical_accuracy: 0.9155  
Epoch 5/10  
1250/1250 [=====] - 80s 64ms/step - loss: 0.2034 -  
sparse_categorical_accuracy: 0.9284  
Epoch 6/10  
1250/1250 [=====] - 78s 62ms/step - loss: 0.1748 -  
sparse_categorical_accuracy: 0.9396  
Epoch 7/10  
1250/1250 [=====] - 78s 62ms/step - loss: 0.1537 -  
sparse_categorical_accuracy: 0.9460  
Epoch 8/10  
1250/1250 [=====] - 78s 63ms/step - loss: 0.1345 -  
sparse_categorical_accuracy: 0.9531  
Epoch 9/10  
1250/1250 [=====] - 78s 63ms/step - loss: 0.1236 -  
sparse_categorical_accuracy: 0.9574  
Epoch 10/10  
1250/1250 [=====] - 79s 63ms/step - loss: 0.1132 -  
sparse_categorical_accuracy: 0.9594
```

1.6 Plotting the Models

```
[15]: # Plotting the 10% Model
```

```
plt.plot(history10.history['loss'])  
plt.title('Model 10% Top-1 Error')
```


```
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()
```



Plotting the 50% Model

```
plt.plot(history50.history['loss'])
plt.title('Model 50% Top-1 Error')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()
```

Plotting the 80% Model

```
plt.plot(history80.history['loss'])
plt.title('Model 80% Top-1 Error')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()
```





2 Homework 2 My Model

2.1 Setting up the Model

```
[16]: # Designing the Model

cifar10_model=tf.keras.models.Sequential()

# Convolutions
↳#####

# First Layer
cifar10_model.add(tf.keras.layers.
↳Conv2D(filters=32,kernel_size=3,padding="same",
activation="relu",
↳input_shape=[32,32,3]))

# Max Pooling Layer
cifar10_model.add(tf.keras.layers.
↳MaxPool2D(pool_size=2,strides=2,padding='valid'))

# Third Layer
cifar10_model.add(tf.keras.layers.
↳Conv2D(filters=64,kernel_size=3,padding="same",
activation="relu"))

# Max Pooling Layer
cifar10_model.add(tf.keras.layers.
↳MaxPool2D(pool_size=2,strides=2,padding='valid'))

# Flattening Layer
cifar10_model.add(tf.keras.layers.Flatten())

# Dropout Layer
cifar10_model.add(tf.keras.layers.Dropout(0.5,noise_shape=None,seed=None))

# Neural Network
↳#####

# Adding the first fully connected layer
cifar10_model.add(tf.keras.layers.Dense(units=128,activation='relu'))

# Output Layer
cifar10_model.add(tf.keras.layers.Dense(units=10,activation='softmax'))
```

```
cifar10_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten (Flatten)	(None, 4096)	0
dropout (Dropout)	(None, 4096)	0
dense (Dense)	(None, 128)	524416
dense_1 (Dense)	(None, 10)	1290

=====
Total params: 545,098
Trainable params: 545,098
Non-trainable params: 0
=====

[17]: *# Setting up the Different Versions*

```
cifar10_model_10 = cifar10_model  
cifar10_model_50 = cifar10_model  
cifar10_model_80 = cifar10_model
```

[18]: *# Compiling Models*

```
cifar10_model_10.compile(loss="sparse_categorical_crossentropy",  
                        optimizer="Adam", metrics=["sparse_categorical_accuracy"])  
  
cifar10_model_50.compile(loss="sparse_categorical_crossentropy",  
                        optimizer="Adam", metrics=["sparse_categorical_accuracy"])  
  
cifar10_model_80.compile(loss="sparse_categorical_crossentropy",  
                        optimizer="Adam", metrics=["sparse_categorical_accuracy"])
```

2.2 Training My Models

[19]: *# Train 10% Model*

```
history10_2 = cifar10_model_10.fit(x_train10,y_train10,epochs=10)
```

Epoch 1/10

157/157 [=====] - 1s 4ms/step - loss: 2.0136 -
sparse_categorical_accuracy: 0.2524

Epoch 2/10

157/157 [=====] - 1s 3ms/step - loss: 1.6405 -
sparse_categorical_accuracy: 0.4162

Epoch 3/10

157/157 [=====] - 1s 3ms/step - loss: 1.4794 -
sparse_categorical_accuracy: 0.4670

Epoch 4/10

157/157 [=====] - 1s 4ms/step - loss: 1.3454 -
sparse_categorical_accuracy: 0.5200

Epoch 5/10

157/157 [=====] - 1s 4ms/step - loss: 1.2822 -
sparse_categorical_accuracy: 0.5402

Epoch 6/10

157/157 [=====] - 1s 4ms/step - loss: 1.1811 -
sparse_categorical_accuracy: 0.5712

Epoch 7/10

157/157 [=====] - 1s 4ms/step - loss: 1.1186 -
sparse_categorical_accuracy: 0.6024

Epoch 8/10

157/157 [=====] - 1s 3ms/step - loss: 1.0437 -
sparse_categorical_accuracy: 0.6332

Epoch 9/10

157/157 [=====] - 1s 4ms/step - loss: 0.9807 -
sparse_categorical_accuracy: 0.6584

Epoch 10/10

157/157 [=====] - 1s 4ms/step - loss: 0.9361 -
sparse_categorical_accuracy: 0.6572

[20]: *# Train 50% Model*

```
history50_2 = cifar10_model_50.fit(x_train50,y_train50,epochs=10)
```

Epoch 1/10

782/782 [=====] - 3s 3ms/step - loss: 1.1954 -
sparse_categorical_accuracy: 0.5849

Epoch 2/10

782/782 [=====] - 3s 3ms/step - loss: 1.0532 -
sparse_categorical_accuracy: 0.6284

Epoch 3/10

```

782/782 [=====] - 3s 3ms/step - loss: 0.9669 -
sparse_categorical_accuracy: 0.6566
Epoch 4/10
782/782 [=====] - 3s 4ms/step - loss: 0.8949 -
sparse_categorical_accuracy: 0.6828
Epoch 5/10
782/782 [=====] - 3s 4ms/step - loss: 0.8353 -
sparse_categorical_accuracy: 0.7051
Epoch 6/10
782/782 [=====] - 3s 4ms/step - loss: 0.7844 -
sparse_categorical_accuracy: 0.7214
Epoch 7/10
782/782 [=====] - 3s 4ms/step - loss: 0.7414 -
sparse_categorical_accuracy: 0.7378
Epoch 8/10
782/782 [=====] - 3s 4ms/step - loss: 0.6896 -
sparse_categorical_accuracy: 0.7550
Epoch 9/10
782/782 [=====] - 3s 4ms/step - loss: 0.6650 -
sparse_categorical_accuracy: 0.7627
Epoch 10/10
782/782 [=====] - 3s 4ms/step - loss: 0.6135 -
sparse_categorical_accuracy: 0.7834

```

[21]: *# Train 10% Model*

```

history80_2 = cifar10_model_80.fit(x_train80,y_train80,epochs=10)

```

```

Epoch 1/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.7775 -
sparse_categorical_accuracy: 0.7328
Epoch 2/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.7122 -
sparse_categorical_accuracy: 0.7521
Epoch 3/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.6692 -
sparse_categorical_accuracy: 0.7646
Epoch 4/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.6280 -
sparse_categorical_accuracy: 0.7780
Epoch 5/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.5989 -
sparse_categorical_accuracy: 0.7871
Epoch 6/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.5764 -
sparse_categorical_accuracy: 0.7969
Epoch 7/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.5442 -

```

```
sparse_categorical_accuracy: 0.8055
Epoch 8/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.5296 -
sparse_categorical_accuracy: 0.8138
Epoch 9/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.5012 -
sparse_categorical_accuracy: 0.8209
Epoch 10/10
1250/1250 [=====] - 5s 4ms/step - loss: 0.4878 -
sparse_categorical_accuracy: 0.8260
```

2.3 Plotting the Loss of My Model

```
[22]: # Plotting the 10% Model

plt.plot(history10_2.history['loss'])
plt.title('Model 10% Top-1 Error')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()

# Plotting the 50% Model

plt.plot(history50_2.history['loss'])
plt.title('Model 50% Top-1 Error')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()

# Plotting the 80% Model

plt.plot(history80_2.history['loss'])
plt.title('Model 80% Top-1 Error')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train'], loc = 'upper left')
plt.show()
```

